Requests for Startups

RFS is our tradition of sharing ideas we'd like to see founders tackle. These represent just a fraction of what we fund — if one excites you, take it as extra validation to dive in, but you don't need to work on these ideas to apply to YC.

Fall 2025 | Summer 2025 | Spring 2025 | Winter 2025 | Summer 2024

Spring 2025

From the AI breakthroughs of the last few months, a wave of new startup opportunities have been unlocked. We used to publish requests for startups no more than once a year, but we decided to publish this list just 3 months after our last one to help point founders to the idea spaces that have just opened up.

A Secure AI App Store

By Garry Tan

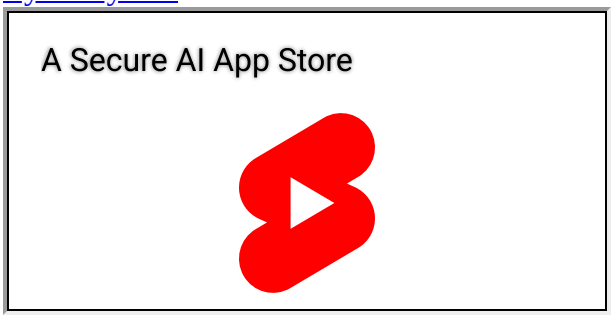

We want a new kind of AI App Store and OS layer that sits on your computer or phone. It should: 1. **Protect User Data** Users control what info each app can see—like your calendar, files, or browsing—but only if the user says so. 2. **Offer One Shared Memory** All your personal details (preferences, past actions, context) stay in this layer, not scattered across a dozen apps. 3. **Help users find the best AI apps** This App Store reviews and vets each AI tool so users can discover and install them safely. 4. **Help Developers Build** Developers get infrastructure that help them avoid reinventing the wheel provided in simple APIs (like computer use, local LLaMA versioning, and app-level access control). 5. **Handle Payments** Make it easy to pay for paid apps or

services. Imagine a travel AI that's great at finding flights that knows you usually travel with your nine-year-old who loves staring out the window, or an AI helper that recommends the original text an idea appeared in when you're reading essays or books. Apps built on this app store would have just a sliver of data you, the user, allows. We need a system like this so AI can be both powerful and private. Some might say the major big tech companies will build this, but there is a time right now where instead, it might be you. Done right, this will create more opportunity for startups and founders: their apps will be smarter with shared memory, and this will turn into a new marketplace that solves both distribution and monetization. If you're building it, please apply. We'd love to help you do it right.

Datacenters

By Diana Hu

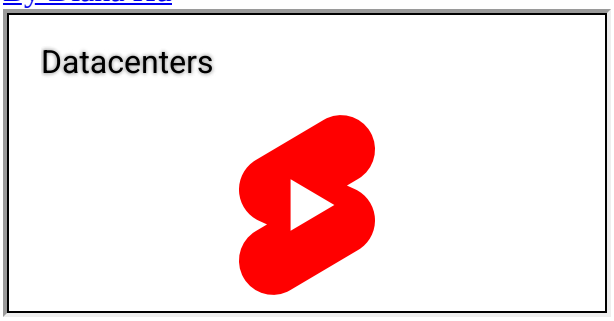

We need more data centers that can be built faster and at a lower cost to support the infrastructure needed for AI progress. Hyperscale datacenter projects take years to complete, and with the current interest and funding, we need new companies and more innovative solutions to accelerate this buildout—whether in power infrastructure, cooling, material procurement, or project management. We can envision what the future will look like: software will handle every aspect of planning and constructing new datacenters or warehouses, from site selection and construction to setup and ongoing management. These applications will be "lights out," with robots operating autonomously 24/7 without human intervention. We want to fund new startups to help create this vision.

Compliance and Audit

By Tom Blomfield

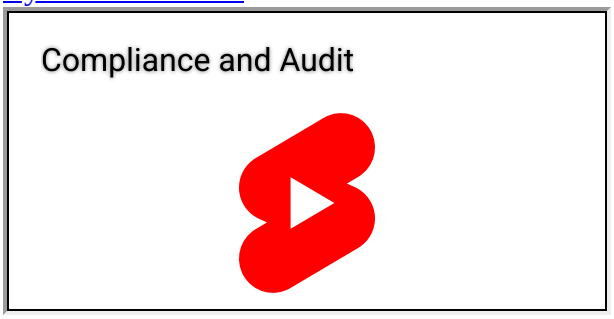

Almost 4 million people in the US and Europe (roughly 1% of the entire workforce) work in compliance and audit. And compliance costs are only rising. From GDPR to Dodd-Frank, from financial AML/KYC to ESG reporting—the regulatory environment keeps expanding. Traditional compliance tasks require reading dense regulations, cross-checking internal policy and procedure documents, manually sampling front-line work done and producing repetitive reports. Auditors often sift through large amounts of unstructured data to detect issues. These manual, time-consuming workflows beg for automation. LLMs excel at this. They can parse regulatory documents, corporate policies, or financial statements and highlight issues, saving human review. These tools will automate much of the testing that auditors currently do by hand: spotting anomalies in data, identifying incomplete records, or highlighting contradictory policies. Instead of sampling a few documents, a well-trained model can look at everything, all at once - "continuously auditing" every company in the world.

Browser & Computer Automation

By Jared Friedman

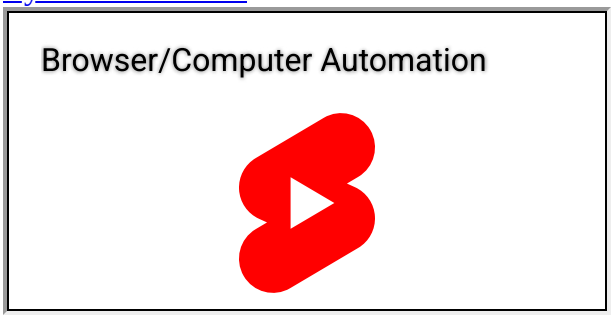

AI agents can now browse the web and use desktop applications. OpenAI's Operator and Anthropic's Computer Use showed this is possible, and there are a number of great open-source options too. Letting AI agents use the web is like taking a brain in a jar and giving it hands. It can now do things. It means that every website and every app now effectively has an API. It means that any workflow that people can do on a computer can be automated. This probably 10x's the addressable use cases for AI agents. We can't wait to see what people build with it.

AI Personal Staff for Everyone

By David Lieb

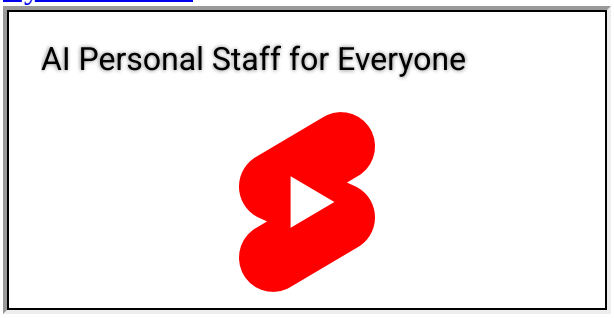

Software is a proven way to bring to everyone what only the rich could afford before. As recently as 2009, only the world's richest people could afford a private driver, but today, Uber, and now Waymo, have brought that to everyone. Another example: when I first started working on what would become Google Photos, I was surprised to learn that rich people hired actual human beings to go through all their photos and edit, label, and organize them. We built AI to automate that for billions of people. Despite the explosion of software in the last decade, wealthy people still employ lots of human staff to provide personal services. These are things like tax accountants, personal lawyers, and money managers, but also personal trainers, private tutors, and even personal doctors. The list goes on and on. Why can only the rich afford this? Because software hasn't been able to replace these types of personalized knowledge work tasks…until now. Over the next few years, we expect AI to get good enough to do most of these jobs. So if you are working to bring a part of this "personal AI staff" to every human on the planet, we'd love to hear from you.

Devtools for AI Agents

By Nicolas Dessaigne

In the past 2 years, we have funded many startups using AI to disrupt legacy players. We are now witnessing the emergence of the next wave: AI agents that go beyond assisting humans and can autonomously make decisions.

With the release of o1, and soon o3, these agents are getting significantly better at reasoning, allowing them to fully replicate and even improve upon tasks performed by humans. AI agents will become ubiquitous across industries and in our daily lives. Imagine a world where every person is empowered by a team of specialized AI agents, working seamlessly in the background to amplify their productivity and creativity. To accelerate this future, we are looking to fund startups building devtools for AI agents. These can include: • Agent builders: Companies enabling their customers to easily create and deploy custom agents, like Wordware (YC S24) or Stack AI (YC W23). • Agent building blocks: Tools, APIs, or platforms that enhance agent capabilities, enabling them to perform more complex actions and achieve greater impact. If you're building in this space, we'd love to hear from you and help you shape the future of software.

The Future of Software Engineering

By Pete Koomen

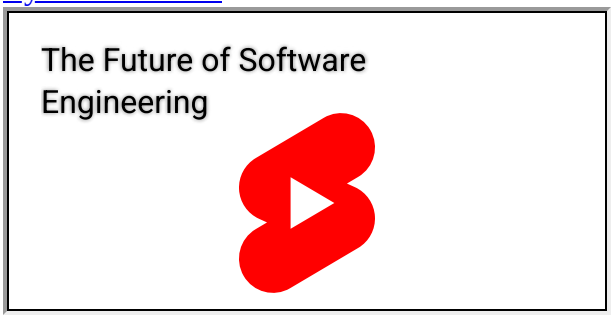

Language models can already write code better than most humans. This is going to bring the cost of building software down to zero. So will agents kill the job of software developer? No! We'll need more human software engineers in the future because software is going to run almost everything. These humans won't write much code directly. Instead, they'll manage teams of agents that build software for them. In addition to writing code, agents will perform most of the other specialized tasks required to build software: QA, deployment, security & compliance audits, translations, operations, etc. We'd like to fund startups that enable small groups of generalist software developers to manage large teams of agents working together to build and ship lots of software. If you're interested in building tools for the future of software engineering, we'd love to hear from you.

AI Commercial Open Source Software (AICOSS)

By Dalton Caldwell

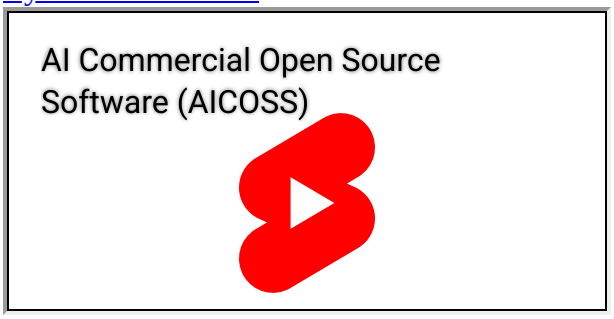

There is a pattern with open source and open source startups. First we had proprietary Unix, then we had Linux, then we had RedHat. First, we had BitKeeper, then we had Git, then we had Github and Gitlab. There is a huge opportunity to build startups that offer support and services to help people use open source AI. It's not uncommon for the organization that releases the open source code to not want to focus on providing commercial support. For instance, Google and Facebook have open sourced a number of tools but aren't always in the business of providing commercial support for businesses to use those tools— which has yielded an opportunity for startups. There are going to be many winners in the space of open source AI, but the launch of DeepSeek should provide ample new ground that a founder might want to cover to help businesses make use of these systems. If you are interested in building in Open Source AI for the enterprise, we would like to hear from you.

AI Coding Agents for Hardware-Optimized Code

By Diana Hu

AI hardware is still constrained by software. Nvidia dominates largely because CUDA's hand-optimized code is used in AI models. Competing hardware—AMD, custom silicon—often underperforms not just because of inferior chips but because writing system-level code (kernels, drivers) is very difficult, and not enough software engineers are working on it. However, now with reasoning models like Deepseek R1 or OpenAI o1 and o3, these could generate hardware-optimized code that rivals—or surpasses—human CUDA code. We'd love to see more founders work on AI-generated kernels that make more hardware alternatives work for AI. This isn't just about performance. It's about breaking dependencies. Founders working on this could reshape the hardware ecosystem.

B2A: Software Where the Customers Will All Be Agents

By Dalton Caldwell

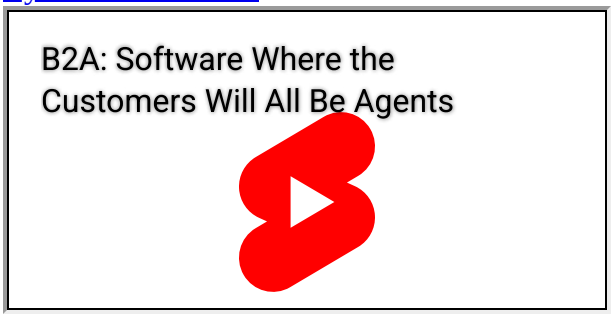

A significant percentage of internet traffic consists of non-humans scraping and looking for information. These programs often end up impersonating humans to fill out forms or look for changes, most people build websites with human users in mind rather than scrapers. With the advent of AI and agents, it seems like it's a good idea to build software and services where agents as customers are actively supported and documented rather than an edge case. For example, APIs to help agents pay for hosting credits, or book travel, or engage into contracts with other parties. In the stock market, it's well understood that humans and programs are trading together, and in the future, this will only increase. If you want to build services specifically aimed at best serving agents, we want to hear from you.

Vertical AI Agents

By Jared Friedman

From 2005-2020, the invention of the interactive web application enabled a massive wave of B2B SaaS companies. Over 100 B2B SaaS unicorns were started, and nearly half of all venture funding went to this category. It seems likely that for the next decade, we will see a parallel wave of companies building vertical AI

agents. What is a vertical AI agent? It's software that's built on top of LLMs that's been carefully tuned to be able to automate some kind of real, important work. In recent batches, we've had YC companies build an AI tax accountant, an AI medical biller, an AI phone support agent, an AI compliance agent, an AI quality assurance tester. Some people previously derided ideas like this as "ChatGPT wrappers" but no one who has tried to deploy a system like this to production thinks so. Building systems like this that work in real-world conditions requires real agentic architectures, integrations with legacy systems, and deep domain understanding. It's hard to get systems like this working right, but once you do, the growth can be phenomenal. The value prop of B2B SaaS was to make human workers incrementally more efficient. The value prop of vertical AI agents is to automate the work entirely. Vertical AI agents that reach human-level performance grow extremely quickly. It's entirely possible this opportunity is big enough to mint another 100 unicorns. For every category with a successful B2B SaaS company, you could imagine an ever larger vertical AI company being built. While many founders are already working on these ideas in the most obvious categories, we think it's still relatively undiscovered compared to the size of the opportunity, and there are many large categories still untouched.

Inference AI Infrastructure in the World of Test-Time Compute

By Diana Hu

Until recently compute spend went into pre-training foundation models. But now with Deepseek R1 and OpenAI o1 and o3, there is a new scaling trend that suggests we'll need far more compute at inference time when AI apps actually use these models. As AI apps 10x or even 100x the number of API calls to complex reasoning models, the infrastructure costs will become a real problem. That's where new startups come in. There's room to rebuild the stack here: better software at inference-layer tooling, cheaper ways to handle GPU workloads, and optimizations that let AI apps scale without bleeding money. It's the kind of unsexy-but-critical problem that often creates big opportunities.

Footer

Y Combinator

Programs

- YC Program
- Startup School
- Work at a Startup
- Co-Founder Matching

Company

- YC Blog
- Contact
- Press

- People
- Careers
- Privacy Policy
- Notice at Collection
- Security
- Terms of Use

Resources

- Startup Directory
- Startup Library
- Investors
- SAFE
- Hacker News
- Launch YC
- YC Deals

Make something people want.

Apply

Twitter Facebook Instagram LinkedIn Youtube